

AI Customer Support Assistant

- Single page HTML only FrontEnd (Chat UI with streaming), and Python Backend
- Backend connects to: SQLite (users, metadata), ChromaDB (vector search)
- Use local LLM (Gemma4 2b on llama.cpp), and OpenRouter Free API as alternative
- Use lightLLM to switch between local and cloud LLMs
- LLM for response generation, Redis for caching frequent queries
- Flow: User asks → Backend retrieves relevant docs from ChromaDB → Augments prompt → Sends to LLM → Streams response back to user with source citations
- Result: Accurate, verifiable, domain-specific answers grounded in company data
- Deploy everything local in one repo (no Docker)